

# Deep Learning for image analysis

## Part II - ConvNet



**JB Fiche**, CBS-Montpellier & Plateforme MARS-MRI

**Francesco Pedaci**, CBS-Montpellier

**Volker Bäcker**, CRBM & MRI

**Cédric Hassen-Khodja**, CRBM & MRI

# Recap - vocabulary



activation function

generalization

overfitting

gradient descent

ground truth

hyperparameter

matplotlib

neuron

one-hot encoding

learning rate

**Backpropagation**

Categorical Cross-Entropy

CNN, ConvNet

Dropout

Keras

Max-Pooling

MNIST

Momentum

**Nonlinearity**

**ReLU**

SGD

softmax

TensorFlow

Vanishing Gradient Problem

VGG16

data augmentation

transfer learning

**epoch**

**weights**

**bias**

hidden layer

**batch size**

**classification**

**regression**

convolution

**loss function**

**testing set**

**validation set**

**training set**

**deep**

optimizer

**fully connected layer**

# Goal of the training :

- Understand what an **Artificial Neural Network (ANN)** is and what are the main parameters to characterize them
- What is a **Convolutional Neural Network (CNN)** and why is it used for image processing
- What are the **fundamentals for building and training a CNN using Keras**
- Understand the **most common applications** and **where to find the tools for your applications**

# Outline :


## **I. Example #4 :**

- A. create an image classifier with a dense network
- B. Introduction to convolutional network

## **II. Exercise #5 : train and optimize your own convNet**

## **III. Real-case examples (Examples # 6 & 7)**

# Working with the MNIST database :



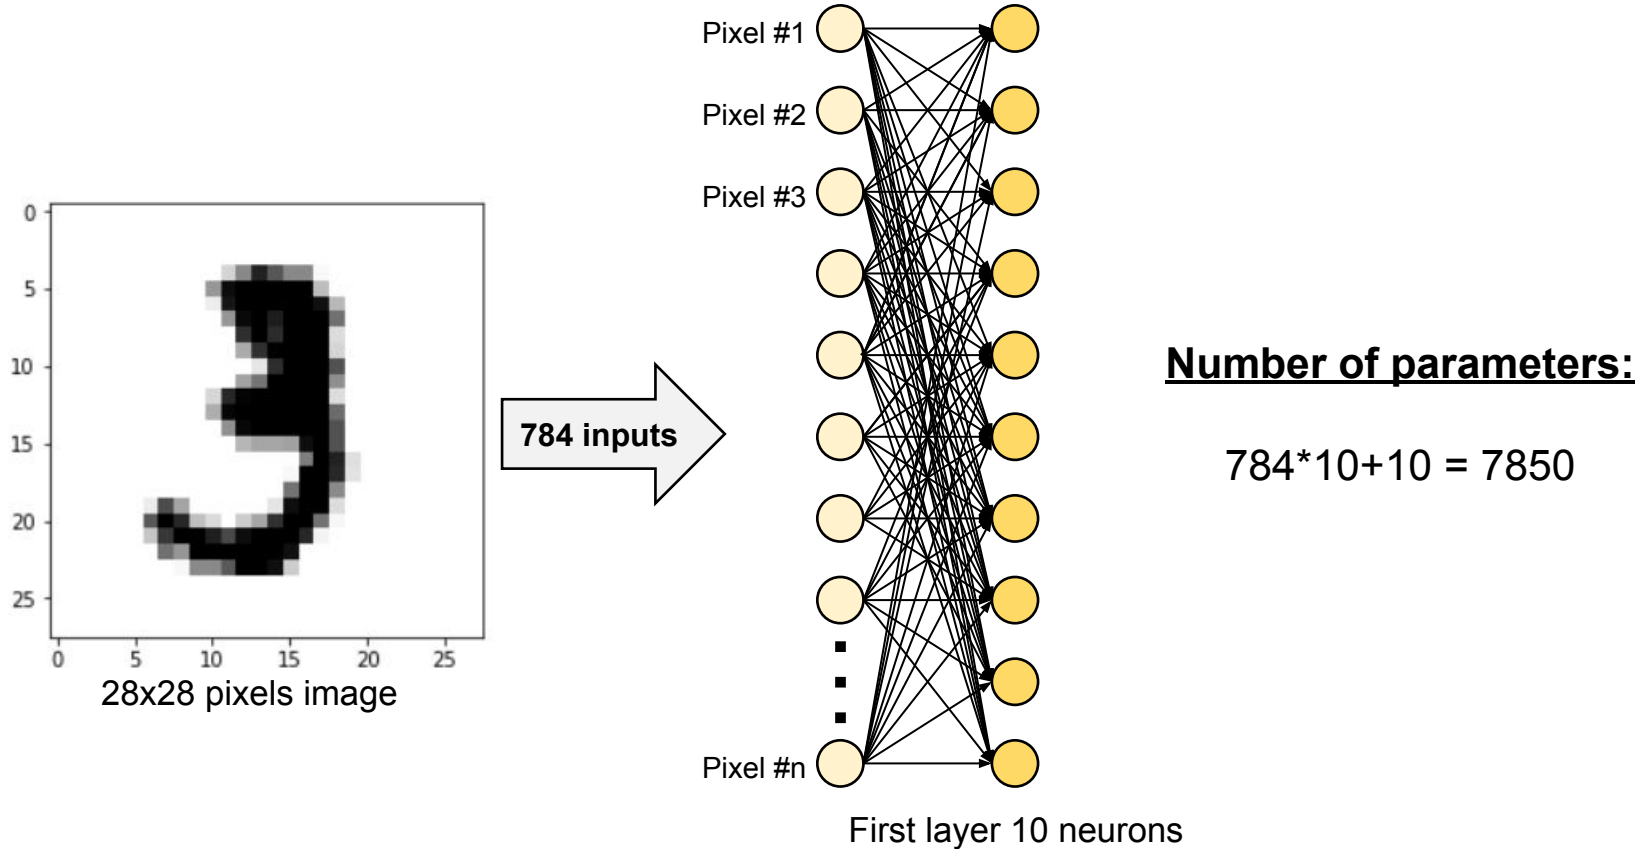
Large database of **handwritten digits** that is commonly used for machine learning.

It contains :

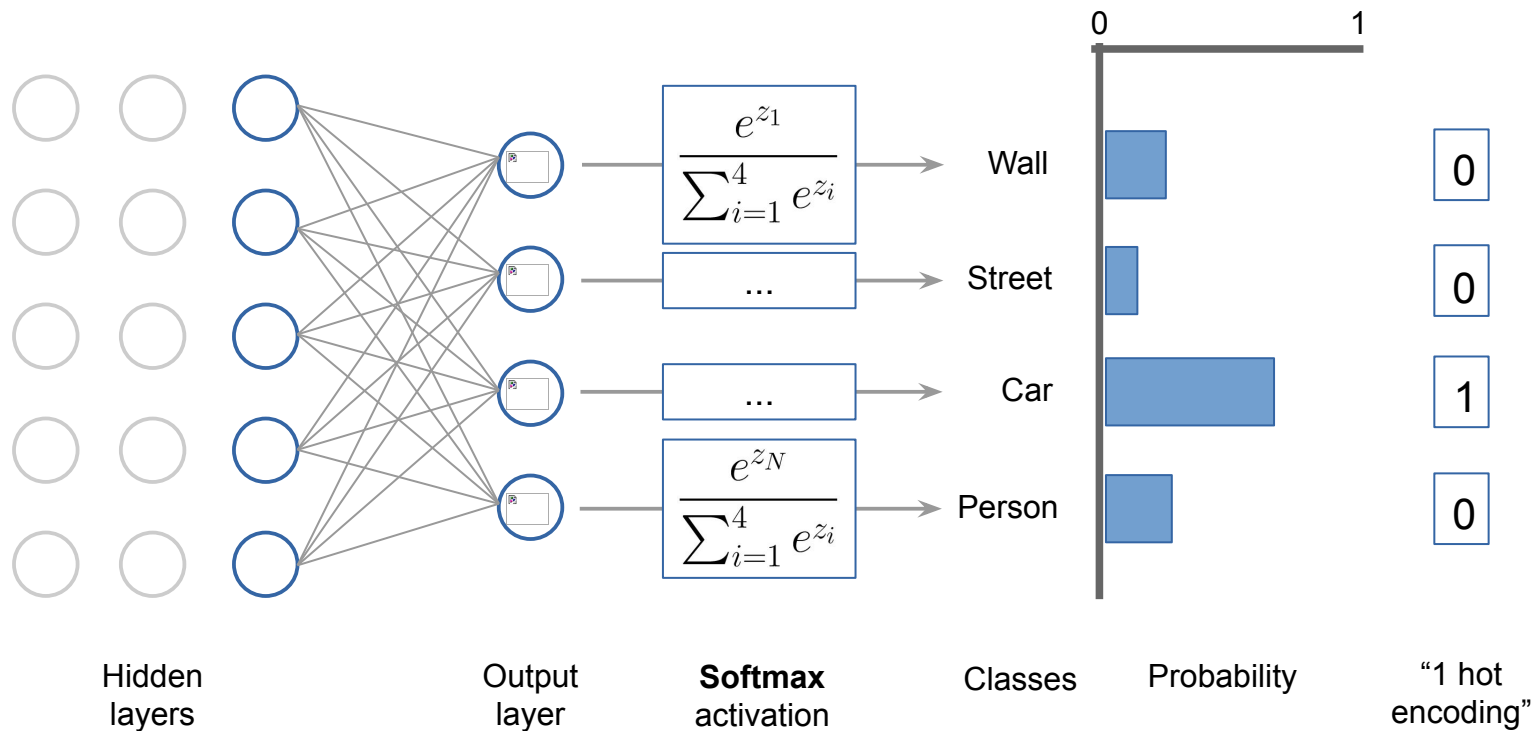
- 60.000 images for training
- 10.000 images for testing/validation

**Ex4\_MNIST\_dense\_vs\_convolutional\_nn.ipynb**

# Image classification with a dense network :





# Classifier with multiple classes : **softmax** activation


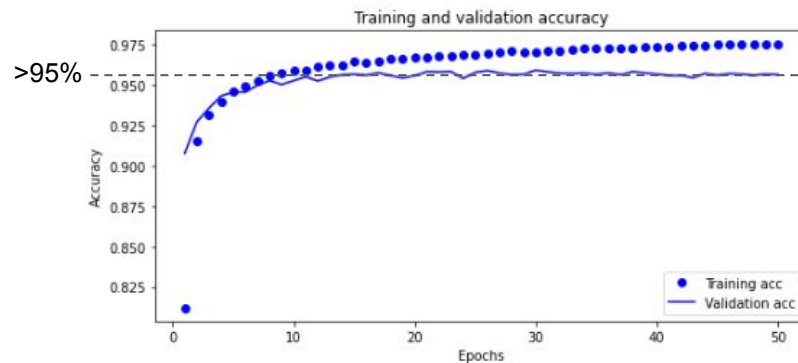


# MNIST classification with a dense network :

2 layers of 10 neurons : **7960 parameters**




3 layers of 15 neurons : **12175 parameters**






# Dense network for large images?



200x200 pixels image



First layer 10 neurons

**Number of parameters:**


$$40.000 \cdot 10 + 10 = 400.010 (!!!)$$

# Image analysis with convolutional network :

Densely connected networks are **not suited** for image analysis:


- Too many parameters, even for small images
- Loses the local information around each pixel
- The network architecture depends on the image size

Go local

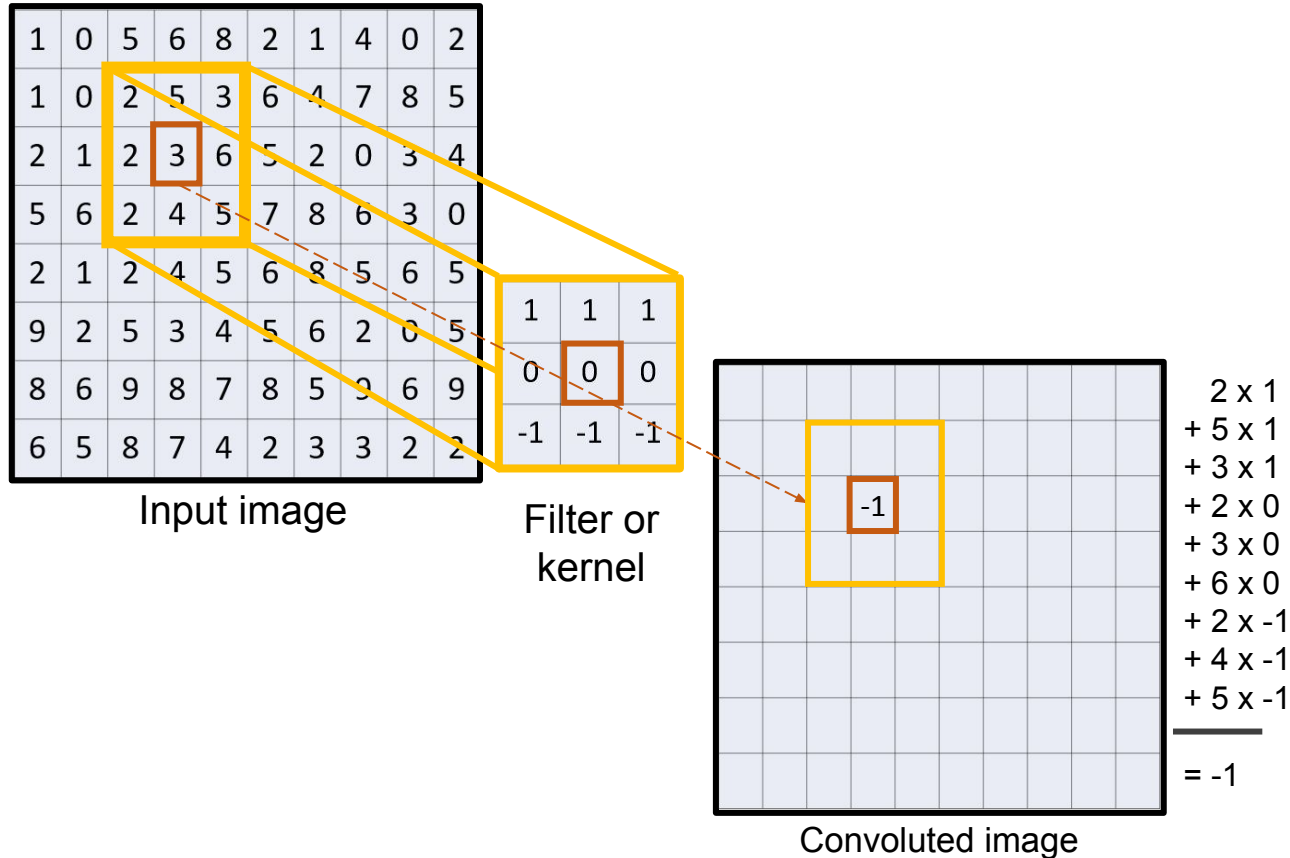


Learn local representations  
with **convolutional** network.


# What is convolution?



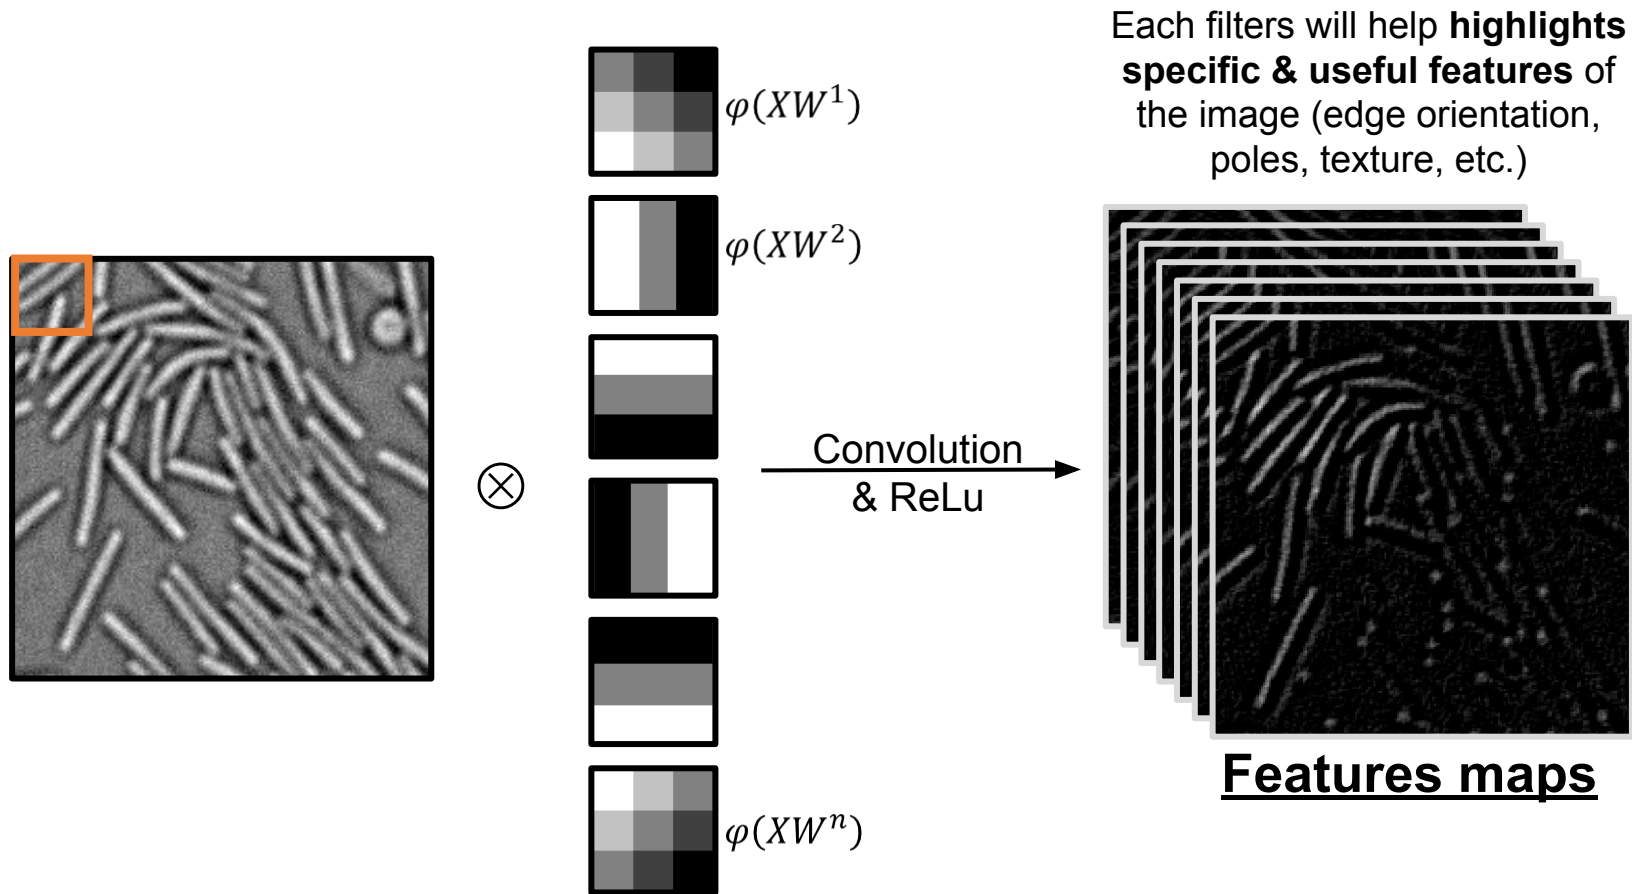
# Convolution operation in one scheme :



# Convolution operation performed by a single neuron:



# Features map :



# ConvNet syntax :

Convolution part

Dense part


```
modelCNN = Sequential([  
    # Convolution Layer 1  
    → Conv2D(16, (3, 3), activation='relu', input_shape=(28, 28, 1)), # 16 different 3x3 kernels -- so 16 feature maps  
    → MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel  
  
    # Convolution Layer 2  
    → Conv2D(16, (3, 3), activation='relu'), # 16 different 3x3 kernels  
    → MaxPooling2D(pool_size=(2, 2)),  
  
    # Convolution Layer 3  
    → Conv2D(16, (3, 3), activation='relu'), # 16 different 3x3 kernels  
  
    Flatten(), # Flatten final 3x3x16 output matrix into a 144-length vector  
  
    # Fully Connected Layer 4  
    → Dense(15), # 15 FCN nodes  
    Activation('relu'),  
    → Dense(10), # Necessary for the last layer since we have 10 classes  
    Activation('softmax'),  
])  
modelCNN.summary()
```

→ Convolution layer

→ MaxPooling layer

→ Dense layer


# Goal of the max-pooling layer :







# Goal of the max-pooling layer :


1. **Reduce the spatial resolution** of the feature maps while keeping only the most relevant information
2. **Lowering memory and computing requirements**
3. Create **translation invariance**



200 x 200 pixels



198 x 198 pixels



99 x 99 pixels

# ConvNet syntax :

Convolution part

Dense part

```
modelCNN = Sequential([  
    # Convolution Layer 1  
    → Conv2D(16, (3, 3), activation='relu', input_shape=(28, 28, 1)), # 16 different 3x3 kernels -- so 16 feature maps  
    → MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel  
  
    # Convolution Layer 2  
    → Conv2D(16, (3, 3), activation='relu'), # 16 different 3x3 kernels  
    → MaxPooling2D(pool_size=(2, 2)),  
  
    # Convolution Layer 3  
    → Conv2D(16, (3, 3), activation='relu'), # 16 different 3x3 kernels  
  
    Flatten(), # Flatten final 3x3x16 output matrix into a 144-length vector  
  
    # Fully Connected Layer 4  
    → Dense(15), # 15 FCN nodes  
    Activation('relu'),  
    → Dense(10), # Necessary for the last layer since we have 10 classes  
    Activation('softmax'),  
])  
modelCNN.summary()
```

→ Convolution layer

→ MaxPooling layer

→ Dense layer

# MNIST classification with a ConvNet :

```
modelCNN = Sequential([  
  
    # Convolution Layer 1  
    Conv2D(16, (3, 3), activation='relu', input_shape=(28, 28, 1)),  
    MaxPooling2D(pool_size=(2, 2)),  
  
    # Convolution Layer 2  
    Conv2D(16, (3, 3), activation='relu'),  
    MaxPooling2D(pool_size=(2, 2)),  
  
    # Convolution Layer 3  
    Conv2D(16, (3, 3), activation='relu'),  
  
    Flatten(),  
  
    # Fully Connected Layer 4  
    Dense(15),  
    Activation('relu'),  
    Dense(10),  
    Activation('softmax'),  
])
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
conv2d_6 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_4 (MaxPooling2D)	(None, 13, 13, 16)	0
conv2d_7 (Conv2D)	(None, 11, 11, 16)	2320
max_pooling2d_5 (MaxPooling2D)	(None, 5, 5, 16)	0
conv2d_8 (Conv2D)	(None, 3, 3, 16)	2320
flatten_2 (Flatten)	(None, 144)	0
dense_4 (Dense)	(None, 15)	2175
activation_4 (Activation)	(None, 15)	0
dense_5 (Dense)	(None, 10)	160
activation_5 (Activation)	(None, 10)	0

Total params: 7,135



Trainable params: 7,135

Non-trainable params: 0

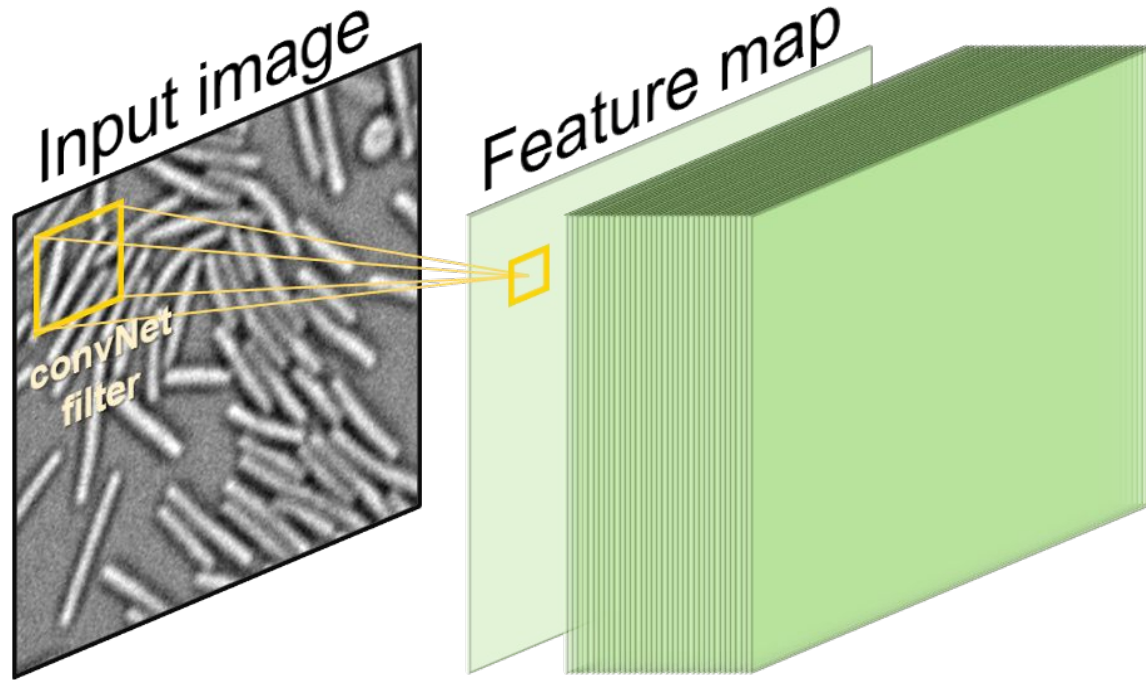
# MNIST classification with a ConvNet :

3 convolution layers of 16 kernel and 2 dense layers of 10 neurons : **7135 parameters**

```
modelCNN = Sequential([  
  
    # Convolution Layer 1  
    Conv2D(16, (3, 3), activation='relu', input_shape=(28, 28, 1)),  
    MaxPooling2D(pool_size=(2, 2)),  
  
    # Convolution Layer 2  
    Conv2D(16, (3, 3), activation='relu'),  
    MaxPooling2D(pool_size=(2, 2)),  
  
    # Convolution Layer 3  
    Conv2D(16, (3, 3), activation='relu'),  
  
    Flatten(),  
  
    # Fully Connected Layer 4  
    Dense(15),  
    Activation('relu'),  
    Dense(10),  
    Activation('softmax'),  
])
```




# ConvNet summary :




Feature maps : there is as **many maps** as **neurons** in the convolution layer

# ConvNet summary :



Feature maps : there is as **many maps** as  
**neurons** in the convolution layer

# ConvNet summary :





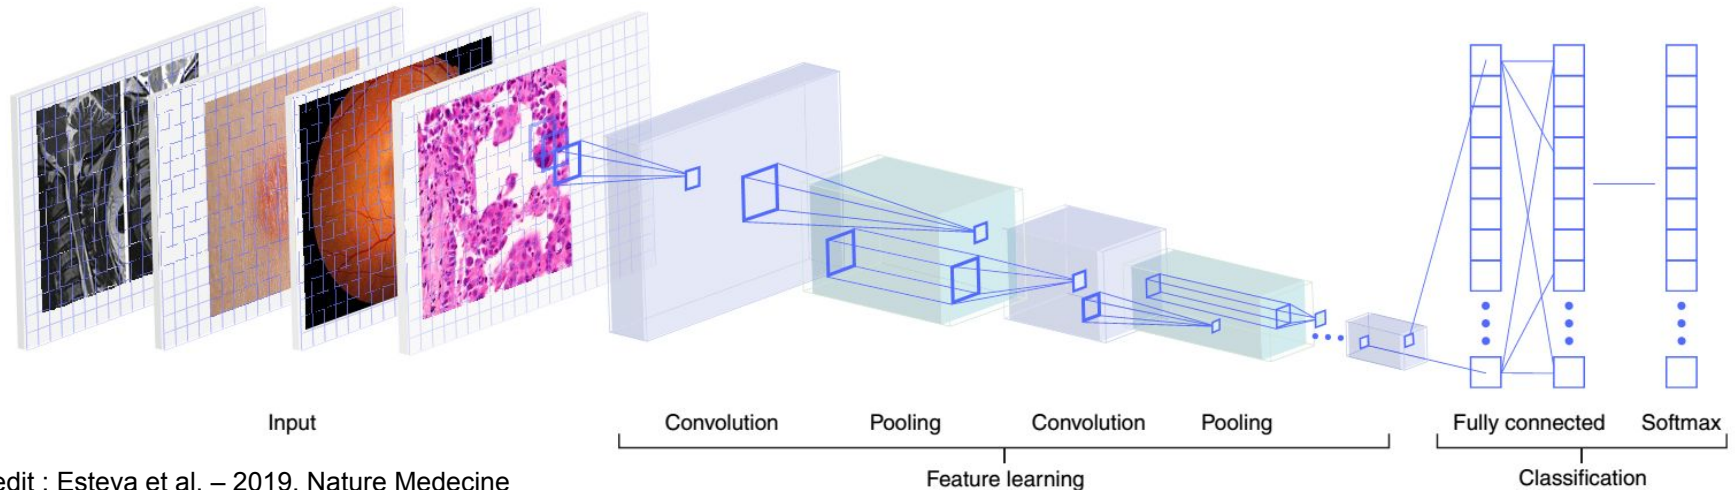
# ConvNet architecture for image classification :

The first part of the network is **using convolution layers to learn the features**

- **Convolution layers** → **features extraction**
- **Pooling** → **downsampling**

The second part is classifying those features using **densely connected layers** in order to predict the right output.

- Lots of parameters → **heavy on the memory**
- Image input size is fixed → **not flexible**





# Outline :

## I. Example #4 :


A. create an image classifier with a dense network

B. Introduction to convolutional network

## II. **Example #5 : train and optimize your own convNet**

## III. Real-case examples

# Optimize your own classifier :



The CIFAR database is a collection of **RGB images** classified into **10 classes** .

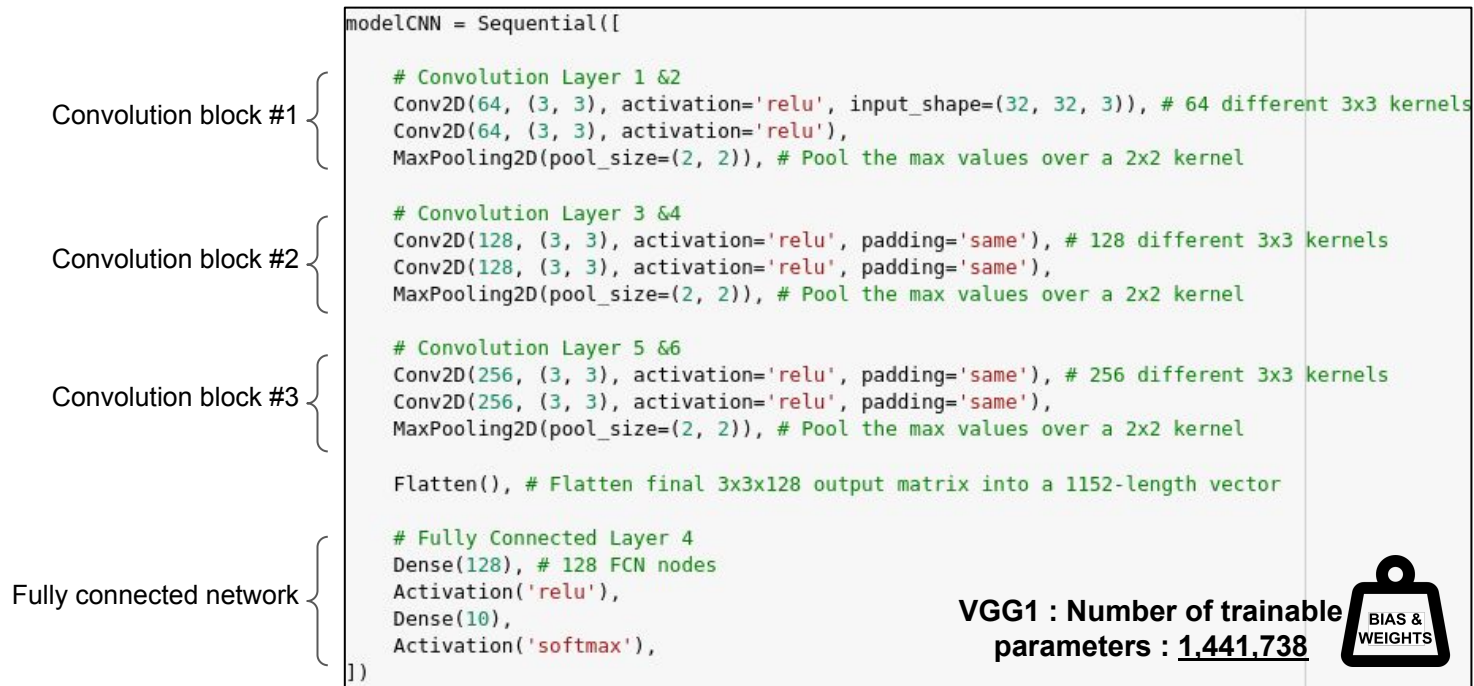
It contains :

- 60.000 images for training. For each class, there are 6000 images.
- 10.000 images for testing/validation


**Ex5\_CIFAR\_convolutional\_nn.ipynb**

# Start with a good baseline model :

A good practice is to **start working with a network architecture that is known to be efficient for your problem**. For example, the VGG16 architecture is easy to implement and well documented for image classification.



# Start with a good baseline model :



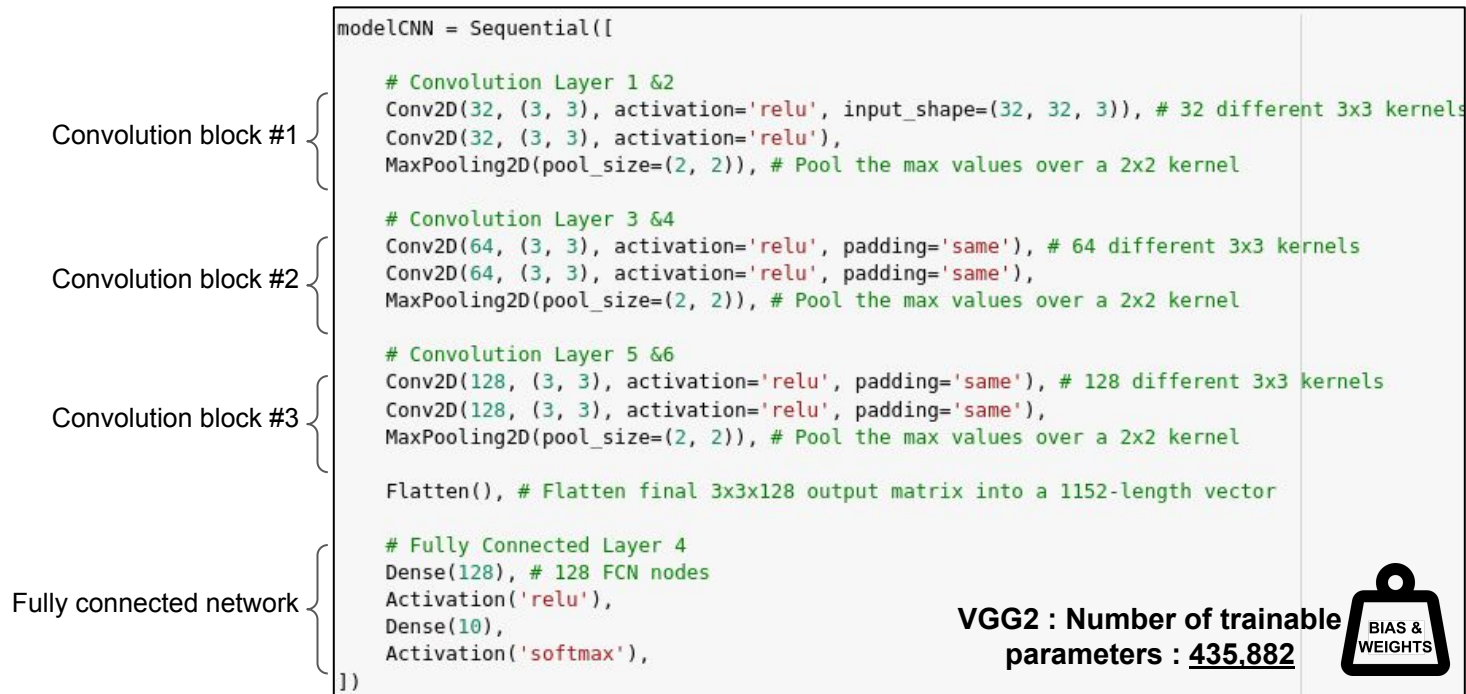
We observe that the training and validation loss & accuracy are diverging after epoch #5. **The network is no longer learning new useful features.**



- Overfitting
- Validation loss & accuracy are noisy
- The global accuracy of the network is ~76.5%



# Reduce the network size :

A good practice is to **start working with a network architecture that is known to be efficient for your problem**. For example, the VGG16 architecture is easy to implement and well documented for image classification.





# Comparison VGG 1 & 2 :

## VGG1




## VGG2




# How to reduce overfitting?

- **Reduce the size of the network** ... but no magical formula to determine how many layers we need
- **Dropout**, randomly “turning-off” neurons of the network



No dropout



Dropout with 40% probability

Dropout is used to avoid co-adaptation of neurons → enforce the fact that neurons should **learn and work independently**.



# VGG baseline and Dropout regularization :

```
modelCNN = Sequential([  
    # Convolution Layer 1 &2  
    Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)), # 32 different 3x3 kernels  
    Conv2D(32, (3, 3), activation='relu'),  
    MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel  
    Dropout(0.2),  
  
    # Convolution Layer 3 &4  
    Conv2D(64, (3, 3), activation='relu', padding='same'), # 64 different 3x3 kernels  
    Conv2D(64, (3, 3), activation='relu', padding='same'),  
    MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel  
    Dropout(0.3),  
  
    # Convolution Layer 5 &6  
    Conv2D(128, (3, 3), activation='relu', padding='same'), # 128 different 3x3 kernels  
    Conv2D(128, (3, 3), activation='relu', padding='same'),  
    MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel  
    Dropout(0.4),  
  
    Flatten(), # Flatten final 3x3x128 output matrix into a 1152-length vector  
  
    # Fully Connected Layer 4  
    Dense(128), # 128 FCN nodes  
    Activation('relu'),  
    Dropout(0.4),  
    Dense(10),  
    Activation('softmax'),  
])
```

Dropout 20% →

Dropout 30% →

Dropout 40% →


Dropout 40% →

Number of trainable  
parameters : 435,882






# VGG baseline and Dropout regularization :




Adding Dropout regularization helps **reducing the overfitting** and slightly **improve the performance of the network (77.7% instead of 75.5%)**



# How to reduce overfitting?

- **Reduce the size of the network** ... but no magical formula to determine how many layers we need
- **Dropout**, randomly “turning-off” neurons of the network
- **Weight regularization  $L_1$  &  $L_2$** , a strategy to force the weights to take only small values during the training



L1/L2 regularization

# VGG baseline and L2 regularization :

```
regularizers.l2(l2=1e-4)
modelCNN = Sequential([

    # Convolution Layer 1 &2
    Conv2D(32, (3, 3), activation='relu', kernel_regularizer='l2', input_shape=(32, 32, 3)), # 32 different 3x3 kernels
    Conv2D(32, (3, 3), activation='relu', kernel_regularizer='l2'),
    MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel

    # Convolution Layer 3 &4
    Conv2D(64, (3, 3), activation='relu', kernel_regularizer='l2', padding='same'), # 64 different 3x3 kernels
    Conv2D(64, (3, 3), activation='relu', kernel_regularizer='l2', padding='same'),
    MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel

    # Convolution Layer 5 &6
    Conv2D(128, (3, 3), activation='relu', kernel_regularizer='l2', padding='same'), # 128 different 3x3 kernels
    Conv2D(128, (3, 3), activation='relu', kernel_regularizer='l2', padding='same'),
    MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel

    Flatten(), # Flatten final 3x3x128 output matrix into a 1152-length vector


    # Fully Connected Layer 4
    Dense(128, kernel_regularizer='l2'), # 128 FCN nodes
    Activation('relu'),
    Dense(10),
    Activation('softmax'),

])
```


Number of trainable  
parameters : 435,882



# VGG baseline and L2 regularization :




Adding L2 regularization helps **reducing the overfitting** but the overall performance of the network is getting worse (61.5% while the baseline was 75.5%).



L2 regularization is not a good method for our classification problem.

# How to reduce overfitting?

- **Reduce the size of the network** ... but no magical formula to determine how many layers we need
- **Dropout**, randomly “turning-off” neurons of the network
- **Weight regularization  $L_1$  &  $L_2$** , a strategy to force the weights to take only small values during the training
- Increase the size of the training set:
  - Add new images to the training set
  - Use **data augmentation**



# VGG baseline and image augmentation :

Create an image generator.  
Important to **select carefully the transformations**. Make sure they are not destroying important information (e.g: object size, shape)



```
# Create the image augmentation generator
# -----

datagen = ImageDataGenerator(width_shift_range=0.1, height_shift_range=0.1, horizontal_flip=True)
it_train = datagen.flow(X_train, Y_train, batch_size=32)
steps = int(X_train.shape[0] / 32)


# Compile the model defining the optimizer and the loss function
# -----

modelCNN.compile(optimizer = 'adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

# Launch the training
# -----

history = modelCNN.fit_generator(it_train,
                                steps_per_epoch=steps,
                                validation_data=(X_val, Y_val),
                                epochs = 50,
                                verbose = 1)
```

# VGG baseline and image augmentation :




Like Dropout regularization, image augmentation helps **reducing the overfitting** and **improve the performance of the network.**

Baseline accuracy : 75.5%



Baseline + Dropout : **77.7% and no overfitting**

Baseline + L2 : 61.5% and no overfitting

Baseline + image augmentation : **82% and less overfitting**



# VGG baseline, dropout & image augmentation :



Like Dropout regularization, image augmentation helps **reducing the overfitting** and **improve the performance of the network**.

Baseline accuracy : 75.5%

Baseline + Dropout : 77.7% and no overfitting

Baseline + L2 : 61.5% and no overfitting

Baseline + image augmentation : 82% and less overfitting

**Baseline + Dropout + image augmentation : 75.5%**



# Adding batch normalization :

```
modelCNN = Sequential([

    # Convolution Layer 1 &2
    Conv2D(32, (3, 3), activation='relu', input_shape=(32, 32, 3)), # 32 different 3x3 kernels
    BatchNormalization(),
    Conv2D(32, (3, 3), activation='relu'),
    BatchNormalization(),
    MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel

    # Convolution Layer 3 &4
    Conv2D(64, (3, 3), activation='relu', padding='same'), # 64 different 3x3 kernels
    BatchNormalization(),
    Conv2D(64, (3, 3), activation='relu', padding='same'),
    BatchNormalization(),
    MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel

    # Convolution Layer 5 &6
    Conv2D(128, (3, 3), activation='relu', padding='same'), # 128 different 3x3 kernels
    BatchNormalization(),
    Conv2D(128, (3, 3), activation='relu', padding='same'),
    BatchNormalization(),
    MaxPooling2D(pool_size=(2, 2)), # Pool the max values over a 2x2 kernel

    Flatten(), # Flatten final 3x3x128 output matrix into a 1152-length vector



    # Fully Connected Layer 4
    Dense(128), # 128 FCN nodes
    Activation('relu'),
    BatchNormalization(),
    Dense(10),
    Activation('softmax'),

])
```

Method introduced in 2015 to **stabilize and speed-up the training** of deep neural networks.

Batch normalization can be implemented during training by calculating the mean and standard deviation of each input variable to a layer per mini-batch and using these statistics to **perform the standardization**.

# VGG baseline, image augmentation & batch normalization:



Like Dropout regularization, image augmentation helps **reducing the overfitting and improve the performance of the network.**

Baseline accuracy : 75.5%

Baseline + Dropout : **77.7% and no overfitting**

Baseline + L2 : 61.5% and no overfitting


Baseline + image augmentation : **82% and less overfitting**

Baseline + Dropout + image augmentation : 75.5%

Baseline + image augmentation + Batch Norm : **89.2%**

# What did we learn?

- How to create and use a **convolutional neural network for image classification** :
  - convolution layer
  - max-pool layer
- Recognize **overfitting** and methods to reduce it while **optimizing the performances of the network** :
  - dropout
  - regularization
  - **image augmentation**
  - batch normalization
  - transfert learning
  - ...



There is no “unique solution”. It strongly depends on your dataset. Need to use a “try & error” strategy.